

flowy



Prompt Like a
Toddler...
Render Like a
God

*Forget arcane syntax:
your dumbest idea just
became a masterpiece.*



Bypass the Valley. Arrive at Skin.

The Common Result




The Goal



Many AI images fall short of realism due to common model biases and a lack of specific control. The difference between a good result and a great one lies in overcoming predictable failure modes like unnaturally smooth surfaces and generic lighting.

Delete Half Your Words. Double Your Realism

Narrative Prompt Structure



The Raw Prompt (Common Mistake)

A photorealistic portrait of a woman in a café. It looks like a photo.

The Structured Prompt (High-Impact)

[Subject]: A photorealistic portrait of a woman, 30s, with a warm, natural smile and freckles.

[Detailed Imagery]: Her hand holds a ceramic cup, soft natural light catches her hair.

[Environment]: A bustling Parisian café with large windows and wooden furniture.

[Mood/Atmosphere]: Cinematic, inviting, golden hour.

[Art Style]: Analog film photography (Kodak Portra 400), shallow depth of field, fine grain.

[Aspect Ratio/Medium]: 4:5 aspect ratio.

The AI Doesn't Know What a Lens Is—But It Knows Leica

Camera & Lens

"Canon 5D Mark IV", "Hasselblad 500CM", "85mm Focal Length".
Simulates specific sensor and lens characteristics.


Key Term

"RAW photo". Instructs the model to generate an image with maximum latent detail and dynamic range, avoiding an "over-processed" look.


Emulate the language of professional photography to force the model to access its most technically precise training data. This moves beyond description into simulation.

Film Stock

"Kodak Portra 400", "CineStill 800T".
Replicates distinct color palettes and grain structures.



Prompt:
"photo.of a woman"



Prompt: "RAW photo of a woman, Canon 5D Mark IV, 85MM FOCAL LENGTH, KODAK PORTRA 400"

Steal Light From 1974

Accurate simulation of light is the cornerstone of realism.

Using terms from 3D rendering forces the model into a physically-based simulation mode.



Volumetric Lighting



Rim Light / Backlight



Chiaroscuro



Global Illumination (GI)

Volumetric Lighting: Creates visible beams of light (god rays) in fog, dust, or mist.

Rim Light / Backlight: Separates the subject from the background with a bright outline, enhancing form.

Global Illumination (GI): Simulates bounced light for softer, more realistic ambient lighting and shadows.

Chiaroscuro: Uses strong contrasts between light and dark for dramatic, high-impact scenes.

Cinematic Fog: Adds atmospheric haze to create depth and enhance volumetric effects.

The Trial of 'Plastic Skin'



The Problem

Models tend to average details in the latent space, resulting in unnaturally smooth surfaces.

The Solution: Forcing High-Frequency Detail

Detailed Skin Texture. This is non-negotiable for realistic portraits. It prompts the model to render fine details like pores, subtle blemishes, and micro-expressions.

Subsurface Scattering (SSS). This is critical. It simulates how light penetrates translucent surfaces like skin, scatters, and exits at a different point, preventing a flat, synthetic look.

Make the AI Stop Blurring Your Vacation Photos

The Bias: SDXL has a strong default bias towards generating images with a shallow depth of field (blurry backgrounds/bokeh), which is often undesirable for landscapes or architectural shots.

SDXL Default Bias



SDXL Default Bias

Overriding the Bias



Overriding the Bias

A Two-Pronged Strategy

Positive Prompts (Demand Sharpness):
Add photographic terms for full-field clarity.
'deep focus', 'long depth of field', 'f/16'.

Negative Prompts (Suppress the Default):
Simultaneously instruct the model what to avoid.
'blurry background', 'bokeh', 'shallow depth of field'.

The AI Listens When You Yell

Tell the AI What You Hate

Keyword Weighting

Assign varying levels of importance to fine-tune the model's focus.

Syntax: Example: (a red car:1.3) on
(keyword:factor) a (blue street:0.8) strongly
emphasizes the car's
redness while
downplaying the street's
blueness.

Alternative: (keyword) increases strength by 1.1x,
[keyword] decreases it by 0.9x.

The Power of Negative Prompts

Your primary tool for refinement. Specify what the model should avoid.

Use Cases: Remove common artifacts, suppress unwanted styles (e.g., cartoon, 3d render, painting), and counter model biases.

A Solid Starter Negative Prompt

ugly, deformed, bad hands, watermark,
signature, text, blurry, bad anatomy, bad
proportions, duplicate, morbid, mutilated,
extra limbs, (worst quality, low quality,
normal quality:1.4)

The Artist's Responsibility: Prompting with Intent

Beyond Technical Bias


Foundational models reflect the biases in their training data. Without specific guidance, they may generate stereotypical or non-representative results.

Your Role

Mastery includes the ability to guide the model toward fair and accurate representations. Be specific and intentional in your descriptions of subjects to ensure diversity.

Ethical Guardrails

These tools can be used to create misleading or harmful content (e.g., "deepfakes"). Professional use requires a commitment to ethical creation, avoiding non-consensual imagery and disinformation.



Let the AI Do Two Jobs, Not One

For the highest quality, SDXL uses a two-stage process. The Base model generates the overall composition and content, and the Refiner adds high-frequency details and polish.



The Golden Ratio

A common strategy is an 80/20 split. For 30 total steps, run 24 steps on the Base model and the final 6 steps on the Refiner.

Advanced Tip: Dual Prompts

Use a detailed content prompt for the Base model and a simpler, quality-focused prompt for the Refiner (e.g., `cinematic, 8k, UHD, high quality, film grain`).

The 'Make It Realer' Button

Standard upscaling can result in soft or artificial-looking images. Advanced workflows use multiple passes and specialized tools to add detail back into the image.



Key Techniques from the Community

Noise Injection: In some workflows (like ComfyUI), strategically adding a small amount of noise during the upscaling process can counteract the model's tendency to smooth surfaces, resulting in more realistic textures.

Specialized Upscaler Models: Don't just use the default. Experiment with community-trained models designed for specific purposes, such as `4x-UltraSharp`, `4xClearRealityV1`, or `1xSkinCompact` for fixing skin.

Iterative Upscaling: A powerful technique involving upscaling an image, downscaling it slightly, and then running it through the full upscaling workflow again to bake in details.

The Extended Arsenal: Specialized Tools (LoRAs)

Base SDXL



Base SDXL can struggle to create true high-contrast or low-key images.

With epi_noiseoffset LoRA



The `epi\_noiseoffset` LoRA grants control over dynamic range.

What is a LoRA?

A Low-Rank Adaptation is a small, specialized model that fine-tunes a specific aspect of the main model without altering it.

Use Case: Fixing Dynamic Range

The `epi\_noiseoffset` LoRA is a well-regarded community tool that solves SDXL's tendency to produce mid-tones, enabling true high-key and low-key photographic styles.

Beyond Utility

LoRAs can also be used to train specific characters, objects, or artistic styles (e.g., a "Chiaroscuro" LoRA to enhance dramatic lighting).

A Gallery of Synthesis

Combining structured prompts, precise modifiers, and advanced workflows unlocks limitless creative potential.



A candid action photograph of a chef tossing flaming ingredients in a hostling kitchen, perfect motion blur, frozen spice particles, volumetric lighting, rich colors, film grain.



Macro photography of a dewdrop on a spiderweb, refraction of a forest incide the drop, extreme detail, shallow depth of field, natural light, photorealistic.



An emetional portrait of a musician in a moment of intense passion, close-up, sweat en brow, closed eyes, natural lighting, detafled skin and hair, candid photography.



A wide-angle architectural photograph of a futuristic llbrary, complex geometric structures, veluxetric sunlight, warm tones, massive scale, detailed textures, atmospheric.



A fantastical photograph of a bio-luminescent jellyfish in a dark underaster cave, electric blue and green light, glewing details, deep black background, ethereal.

The Photorealism Toolkit

Parameter Checklist (SDXL)

Resolution: 1024x1024 (or other native resolutions like 832x1216). Start here before upscaling.

CFG Scale: 7-10 (A balanced range for creativity vs. prompt adherence).

Steps: 25-40 (Good balance of quality and speed). Use an 80/20 split for Base/Refiner.

Sampler: `DDIM` or `Euler a` are great starting points.

Core Modifier Lexicon

Camera: `RAW photo`, `Canon 5D Mark IV`, `85mm Focal Length`

Film: `70mm film`, `Kodak Portra 400`, `Subtle Film Grain`

Lighting: `Volumetric Lighting`, `Subsurface Scattering (SSS)`, `Global Illumination (GI)`, `Rim Light`

Texture: `Detailed Skin Texture`, `Aged Patina`, `Wet Look`

The Path of Iteration

There is no single “**magic prompt.**” True mastery lies in a continuous cycle of iterative experimentation. Use these tools to methodically test, adjust, and build your own intuitive understanding.

Simple inputs can be a starting point. But a detailed prompt, aligned with the model's understanding, is a collaboration. It transforms a request into a fully realized world.

Photorealistic, shot on 70mm IMAX film, an astronaut in a detailed, weathered spacesuit, riding a powerful black horse across the dusty red plains of Mars. The sun is low on the horizon, casting long shadows (Golden Hour). Volumetric dust kicks up from the horse's hooves. Earth is visible as a tiny blue marble in the dark sky. A sense of epic, lonely adventure. hyperrealistic, detailed texture, sharp focus.

